

<p>① Table of Contents</p> <p>Why Parallel Computing? (p. 28/29), Moore's Law (p. 33-44), Challenges / Load Imbalance (p. 48/49) Concepts and Terminology (p. 50-57): Units of measure (FLOP, FLOPs, Bytes; p. 57), von Neumann Architecture (p. 52), Multicore / Multiprocessor / Clusters (p. 54-57) Amdahl's Law (p. 58-78): Load Imbalance example (p. 69), MIPS as metric (p. 73), strong/weak scaling (p. 76-78) Parallel Computing Platforms (p. 79-107): Shared/Distributed memory, communications, synchronization, granularity, overhead (p. 87) Flynn's Taxonomy: SIMD, MIMD, MISD, MIMD (p. 84-94), SPMV / MPMD (p. 95/96), Scalability / Performance (p. 103-107) The Roofline Model (p. 108-146): Estimating nominal performance (p. 120), ridge point (p. 124), Bandwidth / Performance Ceilings, optimization (p. 128-134) Designing Parallel Programs (p. 147-157): Domain / Functional decomposition (p. 150), communication / synchronization (p. 153/152), granularity (p. 156) Memory (p. 158-174): Locality (p. 160/161), Memory Hierarchy / Caches (p. 162-169) Lecture 02: Shared Memory, Concurrency, Threads Processes / Threads (p. 6-21): Fork / Join model (p. 19/20) Memory Model (p. 22-28) Caches (p. 29-44): Types of misses (p. 29), associativity (p. 35-38), coherency (p. 39-44) Mutual Exclusion (p. 46-54): Peterson Lock (p. 54) Threads (p. 55-59): Pthreads (p. 56/57), C++11 Reduction (p. 32-37), Barrier (p. 38), Broadcast (p. 39), Scatter (p. 40), Gather (p. 41), AllGather threads (p. 58/59) OpenMP (p. 60-74): Loop parallelization / compilation (p. 73) Performance Analysis (p. 75-85): Operational Intensity (p. 79-81), Roofline model / Calculation ridge point (p. 82-85) Lecture 03: OpenMP Formal Definition Course / Fine-Grained Parallelism (p. 3) Fork / Join Parallel Regions (p. 3-14): Number of Threads (p. 10), dynamic / static mode (p. 27) Work Sharing Constructs (p. 16-55): Loop Construct (p. 23-26), Sections Construct (p. 27), Single Construct (p. 28), Loop Scheduling (p. 30-32), Nested Loops / Collapsing (p. 33-40), Reductions (p. 41-46), Directive Scoping (p. 47/48), Combined Constructs (p. 49-55) Data Environment (p. 57-64): shared, private, lastprivate, firstprivate Synchronization Constructs (p. 66-71): critical, atomic, mutex, barrier False Sharing (p. 73-81), Library Routines (p. 83-88): Thread IDs, # of processors, time, locks, ... Environment Variables (p. 90-93) Lecture 04: OpenMP UMA / NUMA (p. 3-15): First Touch Policy (p. 8-10) </p>	<p>Lecture 01: Introduction (extended) Introduction (p. 7-30): Processor Binding (p. 17-31): Thread affinity (p. 17), OMP_PLACES Performance (p. 32-37): OMP_WAIT_POLICY (p. 32) Amdahl's Law (p. 39-46): Strong scaling speedup (p. 42) Lecture 04: Linear Algebra / BLAS BLAS (p. 3-27): Levels 1, axpy, gemv, gemm (p. 5/6), Memory Layout (p. 8-13), Performance (p. 79-22) Lecture 05: PCA (notes) 2: Maximum Variance Formulation (p. 2-4), 3: Minimum Error Formulation (p. 4/5), High Dimensional Data (p. 5/6), Kernel PCA / Auto-Associative NN (p. 6/7) Lecture 06: Neural Network (notes) Perceptron (p. 7-107): Shared / Distributed memory, communications, synchronization, granularity, overhead (p. 87) Oja's rule (p. 4/5), Sanger's rule (p. 6) Tutorial 03: Oja's rule: Derivation, Properties (hand-out) Derivation of Oja's Rule (p. 1), Proof of Properties (p. 2-4): w^* is principal component, $\ w^*\ _2 = 1$, w^* maximizes variance at output $E[y^2]$ Lecture 07: ISA / Pipelining Instruction Set Architecture (p. 3-67): Architecture (classes (p. 10), Endianness (p. 20-24), Alignment (p. 25-31), Processor Pipelining (p. 63-80): Throughput (p. 66), CPI (p. 68-70), Speedup (p. 71-74), Hazards (p. 76), Instruction Level Parallelism (p. 82-97): Loop Unrolling (p. 84-89), Code Fusion (p. 90/91) Assembly (p. 93-108) & Lecture 08: MPI Distributed Systems (p. 3-77): MPI Introduction (p. 3-10), MPI-Point-to-Point Communication (p. 18-30): MPI_Send / MPI_Recv (p. 79), data types (p. 20), communication modes (p. 24), buffered send (p. 28) Blocking Collective Communication (p. 37-47): communication modes (p. 42), Alltoall (p. 43), AllReduction (p. 44), Scan / Exscan (p. 45/46) Lecture 09: MPI 2 Unknown size messages / Cycle communication (p. 3-6): MPI_Probe (p. 4), MPI_Sendrecv (p. 6) Eager / Render-Vons protocol (p. 7-10), Non-Blocking Point-to-Point Communication (p. 12-31): MPI_Isend / MPI_Irecv (p. 12-21), MPI_Wait / MPI_Test (p. 15-20), MPI_Waitall / any / some / MPI_Testall / any / some (p. 21/22) Non-Blocking Collective Communication (p. 32-38), Strong / Weak Scaling (p. 40-49): Strong Scaling Efficiency (p. 40), Weak Scaling Efficiency (p. 42), Load Imbalance (p. 48), MPI_Wtime (p. 49) Lecture 10: MPI 3 (1/10, Hybrid MPI + OpenMP) MPI I/O (p. 3-25): File Management (p. 10/11), Data Access (p. 72-74) Hybrid MPI + OpenMP (p. 22-46): Process Mapping (p. 30), Thread Safety (p. 31-34), Ordering Semantics (p. 38-40), MPI_Probe / MPI_Mrecv (p. 42) Lecture 11: Finite Differences / Structured Grids Diffusion (p. 2-30): Random Walk (p. 9-13), Derivation diffusion equation </p>
--	---

(2) (p. 14-19), Fick's 1st/2nd Law (p. 20-26), Diffusion on Infinite Bar (p. 27-29, see also theory) Exam Tips Searching in PDFs:
Diffusion Equation (p. 30) Finite Differences (p. 31-37): FD derivation (p. 31/32), Central/Backward/Forward Difference (p. 33), FD for Diffusion (p. 35-36), Tridiagonal systems/Thomas algorithm (p. 37-39), Von Neumann Stability Analysis (p. 40-43), Modified Equation/Consistency (p. 44-46), Crank-Nicholson Method (p. 47-54), 2D Diffusion (p. 55-57), Peaceman-Rachford Methods / ADI (p. 58-60), 3D Diffusion (p. 61) FFTs + PDEs (p. 63-68) Lecture 11: Particles / N-Body Solvers Particles Introduction (p. 2-52), Cell Lists (p. 53-64), Remeshing (p. 66-75)
Lecture 12: Particle Strength Exchange (Notes) Approximation of a Continuous Function using Particles (p. 1), Smooth Particles (p. 1-5): See Theory for substitution Particle Strength Exchange (p. 5-7): See theory for derivation of (26) and (37) Lecture 13: Data-Level Parallelism / Vectorization ISA Extensions (p. 3-9): Types of Extensions (p. 9) Vector Registers, Alignment, Aliasing (p. 11-27): SIMD Lanes (p. 16), Alignment (p. 27), Aliasing (p. 22-27) Data Oriented Design (p. 28-30): Array of Structures/Structures of Arrays (p. 35-40) Compiler Vectorization/Intel Intrinsics (p. 42-67): Intel Intrinsics (p. 48-67) Intel SPMD Program Compiler (ISPC) (p. 63-89):
Tutorial 01: Scientific Computing Vectors (p. 5) Tutorial 02: BLAS Tutorial 03: Oja's Rule
Taylor Series (p. 7) Tutorial 04: Name Mangling, Debugging, Memory Leaks, Profiling Name Mangling/ Symbol Table (p. 3-7), Debugging (p. 12-16), Memory Leaks (p. 17-20), Profiling (p. 21-24)
Tutorial 05: Parallel Data Storage Tutorial 06: Cell Lists Building Cell Lists (p. 9)
HW 01: Roofline Model / Performance Measures 3.) Operational Intensity: DAXPY / SGEMV / DGEMM, 4.) Diffusion equation 4.) Roofline Model: Calculating Peak Performance, Diffusion Benchmark 5.) Am-dahl's Law HW 02: Brownian Motion / Bug Hunting 1.) Brownian Motion: Reduction (histogram) 2.) OpenMP Bug Hunting HW 03: P(A / Oja's Rule) 1.) Covariance Method: LAPACK 2.) PCA with Oja's Rule: Hebb's Rule, Oja's Rule, Sanger's Rule HW 04: Blocking MPI Routines / Processor Pipelines 1.) 2D Diffusion 2.) MPI Bug Hunt 3.) Processor Pipelining HW 05: Particle Methods / MPI 1.) Roll-up of a Vortex Line: Asynchronous MPI HW 06: 2D Diffusion with Structured Grids and Particles 1.) 2D Diffusion using ADI scheme: ADI discretization 2.) PSE

$$\text{③ Theory Fourier Transform: PDE Solution Fourier Transform: Given } u(x, t), \text{ we have: } \int_{-\infty}^{\infty} u(x, t) e^{-ikx} dx = \sum_{k=-\infty}^{\infty} \hat{u}(k, t) e^{ikx} \quad \text{with } M_k = \int_{-\infty}^{\infty} e^{ikx} \hat{u}(k, t) dx$$

$\hat{u}(k, t) = \int_{-\infty}^{\infty} u(x, t) e^{-ikx} dx$ Inverse Fourier Transform: Given $\hat{S}(k, t)$: $S(x, t) = \frac{1}{2\pi} \int_{-\infty}^{\infty} \hat{S}(k, t) e^{ikx} dk$

Derivatives: Taking the FT of a derivative of order n is the same as multiplication by $(ik)^n$, i.e. $\widehat{u_{xx}}(k, t) = (ik)^2 \hat{u}(k, t)$ **Diffusion on an Infinite Bar:** At $t=0$, $T=Q\delta(x)$ and at $x=\pm\infty$, $T=0$ $\forall t$. Solve $\frac{\partial T}{\partial t} = D \frac{\partial^2 T}{\partial x^2}$. We take the FT and get: $\frac{\partial \hat{T}(k, t)}{\partial t} =$

$D(k)^2 \hat{T}(k, t)$ [derivative property] $\Rightarrow \frac{\partial \hat{T}(k, t)}{\partial t} = -Dk^2 \hat{T}(k, t)$. Integrating w.r.t. t leads to: $\hat{T}(k, t) = f(k) e^{-Dt k^2}$. At $t=0$, we have $\hat{T}(k, 0) = f(k)$ and from the initial cond.

ition $\hat{T}(k, 0) = \int_{-\infty}^{+\infty} Q\delta(x) e^{-ixk} dx = Q$. Combining both leads to $f(k) = Q$ and therefore $\hat{T}(k, t) = Q e^{-Dt k^2}$. We now take the IFT: $T(x, t) = \frac{1}{2\pi} \int_{-\infty}^{\infty} \hat{T}(k, t) e^{ikx} dk = \frac{Q}{2\pi} \int_{-\infty}^{\infty} e^{-Dt k^2} e^{ikx} dk = \frac{Q}{2\pi} e^{-\frac{x^2}{4Dt}} \int_{-\infty}^{\infty} e^{-Dt(k^2 - \frac{x^2}{4Dt})} dk =$

The integral is a Gaussian with $2\sigma^2 = \frac{1}{Dt} \rightarrow \sigma = \frac{1}{\sqrt{2Dt}}$. The general solution to a Gaussian: $\int_{-\infty}^{\infty} e^{-\frac{x^2}{2\sigma^2}} dx = \sqrt{\pi} \frac{1}{\sigma \sqrt{2\pi}}$ gives: $\frac{Q}{2\pi} e^{-\frac{x^2}{4Dt}} = \frac{Q}{\sqrt{4\pi Dt}} e^{-\frac{x^2}{4Dt}}$ **Von Neumann**

Stability Analysis: Stability Analysis of Diffusion Equation: We assume periodic BCs, PDEs with linear, constant coefficients and a uniformly spaced grid. Then, we can write a solution

in the form: $u_j^{(n)} = p^n e^{ikx_j}$. Applying the ansatz to the discretization $u_j^{(n+1)} = u_j^{(n)} + \frac{r\delta t}{\delta x^2} (u_{j+1}^{(n)} - 2u_j^{(n)} + u_{j-1}^{(n)})$ gives: $p^{n+1} e^{ikx_j} = p^n e^{ikx_j} + \frac{r\delta t}{\delta x^2} p^n (e^{ikx_{j+1}} - 2e^{ikx_j} + e^{ikx_{j-1}})$

We introduce $x_{j-1} = x_j - \delta x$, $x_{j+1} = x_j + \delta x$ and write: $p^{n+1} e^{ikx_j} = p^n e^{ikx_j} + \frac{r\delta t}{\delta x^2} (e^{ik(x_j + \delta x)} - 2e^{ikx_j} + e^{ik(x_j - \delta x)})$, which leads to $p = 1 + \frac{r\delta t}{\delta x^2} (e^{ik\delta x} - 2 + e^{-ik\delta x}) =$

$1 + \frac{r\delta t}{\delta x^2} (\cos(k\delta x) + i\sin(k\delta x) - 2 + \cos(-k\delta x) + i\sin(-k\delta x))$. Using $\sin(-x) = -\sin(x)$, $\cos(-x) = \cos(x)$: $p = 1 + \frac{r\delta t}{\delta x^2} (2\cos(k\delta x) - 2)$. For stability, we need to have $|p| \leq 1$, i.e.

$-1 \leq 1 + \frac{r\delta t}{\delta x^2} (2\cos(k\delta x) - 2) \leq 1$. $p \leq 1$ is always fulfilled, so we have: $p \geq -1 \Leftrightarrow \frac{r\delta t}{\delta x^2} (2\cos(k\delta x) - 2) \geq -2$. In the worst case, $\cos(k\delta x) = -1 \Rightarrow -\frac{4r\delta t}{\delta x^2} \geq -2 \Leftrightarrow$

$\frac{r\delta t}{\delta x^2} \leq \frac{1}{2} \Leftrightarrow \delta t \leq \frac{\delta x^2}{2r}$ **FFT+PDEs** For periodic boundary conditions. Idea: Representing Fourier Series, form derivatives by multiplication, solve for next time step, take IFT.

Particle Strength Exchange Smooth Particles Derivation: From (72) to (75): Let $z = \frac{x-y}{\varepsilon}$, $\frac{dz}{dy} = -\frac{1}{\varepsilon} \rightarrow dz = -\frac{dy}{\varepsilon}$. Substituting z into (72) gives: $\sum_{k=-\infty}^{\infty} \frac{f^{(kn)}(x)}{k!} \frac{\varepsilon^k}{\varepsilon^k} \cdot$

$$\int_{-\infty}^{\infty} (-2\varepsilon)^k \frac{1}{k!} \eta(z) \varepsilon dz = \sum_{k=-\infty}^{\infty} \frac{f^{(kn)}(x)}{k!} \int_{-\infty}^{\infty} (-2)^k \varepsilon^k \eta(z) dz = \sum_{k=-\infty}^{\infty} \frac{f^{(kn)}(x)}{k!} (-1)^k \varepsilon^k M_k$$

$$\text{with } M_k = \int_{-\infty}^{\infty} z^k \eta(z) dz \quad \text{PSE Derivation: To get (26): Apply (27) to (24):}$$

$$\frac{\partial F(x_i, t)}{\partial t} = \frac{\partial}{\partial t} \sum_{p=1}^N [f(x_p, t) - f(x_i, t)] V_p \eta_\varepsilon(x_i - x_p) \quad \text{Apply (25): } \frac{dV_i}{dt} =$$

$$\frac{D}{\varepsilon^2} \sum_{p=1}^N \left[\frac{\partial \eta_\varepsilon}{\partial x_p} - \frac{V_i}{V_p} \right] V_p \eta_\varepsilon(x_i - x_p) \rightarrow \frac{dV_i}{dt} = \frac{D}{\varepsilon^2} \sum_{p=1}^N [Q_{ip} V_i - Q_i V_p] \eta_\varepsilon(x_i - x_p)$$

(multiplication of both sides with V_i , terms on RHS with V_p). **Regular Lattice:**

The moments for the given kernel are: $M_1 = (-1) \cdot 1 + 0 \cdot 2 + 1 \cdot 1 = 0$, $M_2 = (-1)^2 \cdot 1 + 0 \cdot 2 + 1 \cdot 1 = 2$, $M_3 = (-1)^3 \cdot 1 + 0 \cdot 2 + 1^3 \cdot 1 = 0$. We get eq. (37) by: Use (26) with (7) and $d=1$:

$$\frac{dV_i}{dt} = \frac{D}{\varepsilon^2} \sum_{p=1}^N [Q_{ip} V_i - Q_i V_p] \frac{1}{\varepsilon} \eta\left(\frac{x_i - x_p}{\varepsilon}\right) \quad \text{We have } V_p = V_i = h, \varepsilon = h, x_i + x_p = (i-p)h$$

$$\text{and therefore: } \frac{D}{h^2} \sum_{p=1}^N [Q_{ip} - Q_i] \eta\left(\frac{(i-p)h}{h}\right) = \frac{D}{h^2} \sum_{p=1}^N [Q_{ip} - Q_i] \eta(i-p) \quad \eta(i-p) \text{ is } 0,$$

$$\text{unless } p=i-1, i, i+1; \text{ so: } \frac{dV_i}{dt} = \frac{D}{h^2} [(Q_{i-1} - Q_i) \cdot 1 + (Q_i - Q_{i+1}) \cdot 2 + (Q_{i+2} - Q_i) \cdot 1] = \frac{D}{h^2} [Q_{i-1} -$$

$2Q_i + Q_{i+1}]$ **Advantages / Disadvantages Particle Methods** Advantages: 1.) conservative 2.) can solve some advection problems exactly (e.g. linear convection) 3.) Lower / No numerical diffusion 4.)

Easier Refinement Disadvantages: 1.) More complex algorithms for finding neighbor particles

2.) Fewer choices of methods to solve linear systems 3.) Harder to impose boundary

conditions **PDE Notation** $\nabla = \left(\frac{\partial}{\partial x_1}, \dots, \frac{\partial}{\partial x_n} \right)^T$, $u_{tt} = \frac{\partial^2}{\partial t^2} u(t, \dots)$, $u_{tx} = \frac{\partial^2}{\partial t \partial x} u(t, x, \dots)$, $\Delta = \nabla \cdot \nabla =$

$\sum_{i=1}^n \frac{\partial^2}{\partial x_i^2}$ **Trigonometric Identities** $e^{ix} = \cos(x) + i\sin(x)$, $e^{i\theta} + 1 = 0$, $\cos(x) = \frac{e^{ix} + e^{-ix}}{2}$, $\sin(x) =$

$\frac{e^{ix} - e^{-ix}}{2i}$, $\sin(-\theta) = -\sin(\theta)$, $\cos(-\theta) = \cos(\theta)$, $\sin(\theta \pm \frac{\pi}{2}) = \pm \cos(\theta)$, $\cos(\theta \pm \frac{\pi}{2}) = \mp \sin(\theta)$,

$\sin(\theta + \pi) = -\sin(\theta)$, $\cos(\theta + \pi) = -\cos(\theta)$, $\pi = 180^\circ$, $e^{-ix} = \cos(x) - i\sin(x)$, $\cos^2(x) + \sin^2(x) = 1$

Amdahl's Law: Speedup $S_n = \frac{1}{1-p+\frac{p}{n}}$, where p =parallel fraction. If $n \rightarrow \infty$, $S_n = \frac{1}{1-p}$

Strong Scaling Analysis We keep the problem size fixed and report $S_p = \frac{t_1}{t_p}$. We plot

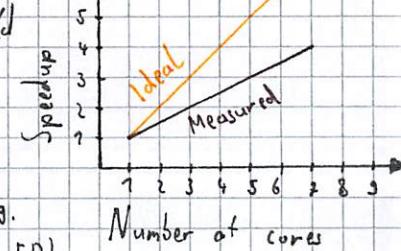
the number of cores/thread (x-axis) against the speedup:

Strong scaling efficiency is defined as $\eta_S(p) = \frac{S_p}{p}$. Would

be 1 in the ideal case. **Weak Scaling Analysis** We in-

crease the number of nodes and the problem size s.t. the work per node is constant. If amount of work is e.g.

(linear with problem size (e.g. 7D diffusion with forward FD),



④ multiplying nodes by c requires multiplying problem size by c . If amount of work is quadratic (e.g. naive N body interactions), multiplication by \sqrt{c} is required.

Weak scaling efficiency: $\eta_w(p) = \frac{T(1)}{T(p)}$. We plot the number of cores/threads against the weak scaling efficiency.

Roofline Model Operational Intensity: $\frac{\text{FLOPs}}{\text{Byte}}$

Given nominal peak bandwidth B (in GB/s) and

nominal peak performance \bar{f}_C (in GFlops/s), the performance of code with operational intensity I is:

$$P = \min(BI, \bar{f}_C). \text{ The ridge point is } I_b = \frac{\bar{f}_C}{B}, \text{ if } I < I_b$$

memory bound and if $I > I_b$ compute bound. Roofline graph: Horizontal line at \bar{f}_C until ridge point, then diagonal line that goes through B at $1/10^\circ$ (lecture 02, p. 85).

Data Types Single Precision / float: 4 bytes, Double Precision / double: 8 bytes.

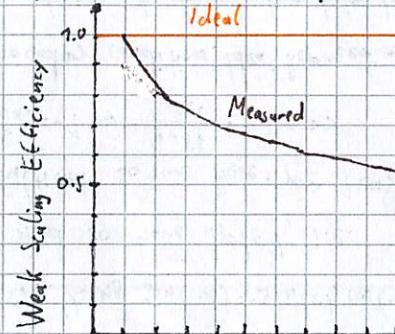
(unsigned) short int: 2 bytes, (unsigned) int / (unsigned) long int: 4 bytes, (unsigned) long

(long int = 8 bytes) 1D Diffusion: Finite Differences Explicit: $u_i^{n+1} = u_i^n + \nu \frac{\Delta t}{\Delta x^2} (u_{i+1}^n + u_{i-1}^n - 2u_i^n)$

Implicit: $u_i^{n+1} = u_i^n + \nu \frac{\Delta t}{\Delta x^2} (u_{i+1}^{n+1} + u_{i-1}^{n+1} - 2u_i^{n+1})$ 2D Diffusion: Finite Differences

$$p_{i,j}^{n+1} = p_{i,j}^n + \Delta t D \left(\frac{p_{i,j-1}^n - 2p_{i,j}^n + p_{i,j+1}^n}{\Delta x^2} + \frac{p_{i-1,j}^n - 2p_{i,j}^n + p_{i+1,j}^n}{\Delta x^2} \right). \text{ With } \alpha_x = \alpha_y = \Delta x \text{ becomes:}$$

$$p_{i,j}^{n+1} = p_{i,j}^n + \frac{\Delta t D}{\Delta x^2} \left(p_{i,j-1}^{n+1} - 2p_{i,j}^{n+1} + p_{i,j+1}^{n+1} + p_{i-1,j}^{n+1} - 2p_{i,j}^{n+1} + p_{i+1,j}^{n+1} \right)$$



⑤ Programming C++ Arrays	SSE Example (SSE)	ISPC Example:
Static creation: double data[10];, Dynamic: double* data = new double[10];, Initialization: int x[] = {1, 2, 3};, Iteration: for(int i=0; i<len; i++) { arr[i]}; Vectors: Required header: #include <vector>, Creation: std::vector<int> vec1 or with length n, initialized to 0: std::vector(n, 0);, Appending: vec.push_back(i);, Length: float *x, *y;	#include <std/stl.h> #include <x86intrin.h>	uniform float const uniform x, uniform float* const uniform y,
vec.size();, Iterating: for(int i=0; i<vec.size(); i++) { vec[i]};, Deleting by index i: vec.erase(vec.begin() + i), Accessing as Array: int* a = &vec[0];, Copying: When initializing new vector: std::vector<int> vec2(vec); otherwise: vec2 = vec1 (assignment copies content), Swapping (constant complexity): vec.swap(vec2);, Swapping for arrays: std::swap();	posix_memalign((void**)&x, 16, N*sizeof(float)); posix_memalign((void**)&y, 16, N*sizeof(float)); const int SIMD_WIDTH = 16/sizeof(float); for(int i=0; i<N; i+=SIMD_WIDTH) {	const uniform int N { for(int i=0; i<N; i+=programCount) { y[i] = x[i] + y[i]; }
Emitting Assembly: g++ -S test.cpp printf: Required header: #include <csdio.h>, Syntax: printf("format string", arg1, arg2, ...);, Format specifiers: %d: Signed int, %u: Unsigned int, %c: Character, %s: String, %p: Pointer addressers Makefiles: E.g.: exec-omp.c g++ -O3 -std=c++11 -Wall -fopenmp main.cpp -o J @ Debugging Shared Memory: g++ -g -fopenmp -fsanitize=thread will report data races. cout: Required header: #include <iostream>, Syntax: std::cout << "Output" << std::endl; Structs: E.g.: struct Student { char name[20]; int id; int age; }, Creating an instance: Student s; Accessing fields: s.id=4;	const _mm ₁₂₈ &x4 = _mm_load_ps(x+i); const _mm ₁₂₈ &y4 = _mm_load_ps(y+i); _mm_store_ps(y+i, _mm_add_ps(x4, y4)); }	foreach (i=0 ... N) y[i] = x[i] + y[i]; }
	free(x); free(y);	AVX with doubles: const _mm ₂₅₆ &x4 = _mm256_load_pd(x+i);
		HW code (cont.) periodic_dist: double dist=x1-x2; if(dist < -5*DOMAIN_SIZE) { return dist+DOMAIN_SIZE; } if(dist > 5*DOMAIN_SIZE) { return dist; } return dist-DOMAIN_SIZE;

⑥ OpenMP Compilation: `g++ -fopenmp` Running: Programs can be run normally, can use MPI Compilation: `mpic++ main.cpp -o main` Running: `mpiexec -n 4 main` Header: `#include <mpi.h>`

OMP_NUM_THREADS to control number of threads, i.e. `OMP_NUM_THREADS=4 ./program` for one run, export `OMP_NUM_THREADS=4` for persistence. Header: `#include <omp.h>`

Only required for library calls. Parallel Regions: `#pragma omp parallel [clause[;] clause...]` with clause: 1.) if([parallel]): scalar-expression → Execute only if true 2.) num_threads[n] → Number of threads in team 3.) default(shared|none) → default data-sharing 4.) private(list) → Uninitialized private copy 5.) firstprivate(list) → Initialized (to master value) private copy 6.) shared(list) 7.) proc_bind(master|close|spread) → Binding, see lecture 04, p. 78

Loop Construct: With implicit barrier: `#pragma omp for [clause[;] clause...]` with clause: 1.) private(list) 2.) firstprivate(list) 3.) lastprivate(list) → Uninitialized private copy, in the end updated according to last recvbuf, int count, MPI_Datatype datatype, MPI_Op op, int root, MPI_Comm comm with

(lexical section 4.) schedule(static [, chunk] | dynamic [, chunk] | guided [, chunk] | runtime | auto) → MPI_OPE{MPI_MAX, MPI_MIN, MPI_SUM, MPI_PROD, ...} (lecture 08, p. 34)

lecture 03, p. 31 5.) collapse(n) → Collapse n (perfectly nested) levels 6.) reduction(red_id: list) → To-Point: `MPI_Isend(const void *buf, int count, MPI_Datatype datatype, int dest, int tag, MPI_Comm comm, MPI_Request *request)` / `MPI_Irecv(void *buf, int count, MPI_Datatype datatype, int source, int tag, MPI_Comm comm, MPI_Request *request)`

7.) nowait Other Programs: * = with barrier. `#pragma omp ...` int tag, MPI_Comm comm, MPI_Request *request); Operations on request: `MPI_Wait(MPI_Request request, MPI_Status *status)`; Wait for completion, `MPI_Test(MPI_Request request, int *flag, MPI_Status *status)`; → flag=1 if completed.

1.) master → Only performed by master thread 2.) single*: By one thread 3.) critical [(name)] → (critical section 4.) barrier* 5.) atomic → Replaced by atomic H/W operations 6.) sections*: Each section assigned to (one) different thread, works also with more sections than threads

Library Routines: 1.) `omp_get_thread_num()` → Get thread 11) 2.) `omp_get_num_threads()` → Threads in team 3.) `omp_set_num_threads(n)` → Set number of threads 4.) `omp_get_wtime()` → Per thread wall time (delta=
 +2.171) Environment Variables: 1.) `OMP_DYNAMIC="TRUE" or "FALSE"` → Dynamic mode, i.e. adjustment of number of threads 2.) `OMP_NESTED="TRUE" or "FALSE"` 3.) `OMP_PROC_BIND="TRUE" or "FALSE"` → Whether threads are allowed to be moved ("FALSE") between processes

Basic Program: `int rank, size; MPI_Init(&argc, &argv); MPI_Comm_rank(MPI_COMM_WORLD, &rank); MPI_Comm_size(MPI_COMM_WORLD, &size); MPI_Finalize();` Blocking MPI Point-to-Point: `MPI_Send(const void *buf, int count, MPI_Datatype datatype, int dest, int tag, MPI_Comm comm); / MPI_Recv(void *buf, int count, MPI_Datatype datatype, int source, int tag, MPI_Comm comm, MPI_Status *status)`
 Tag can be `MPI_ANY_TAG`, source `MPI_ANY_SOURCE` and status `MPI_STATUS_IGNORE`

MPI-Datatype: 1.) char: `MPI_CHAR` 2.) int: `MPI_INT` 3.) float: `MPI_FLOAT` 4.) double: `MPI_DOUBLE` (lecture 08, p. 20)

Reduction: `MPI_Reduce(const void *sendbuf, void *recvbuf, int count, MPI_Datatype datatype, MPI_Op op, int root, MPI_Comm comm)` with

MPI_OPE{MPI_MAX, MPI_MIN, MPI_SUM, MPI_PROD, ...} (lecture 08, p. 34)

Non-Blocking Point-to-Point: `MPI_Isend(const void *buf, int count, MPI_Datatype datatype, int dest, int tag, MPI_Comm comm, MPI_Request *request)` / `MPI_Irecv(void *buf, int count, MPI_Datatype datatype, int source, int tag, MPI_Comm comm, MPI_Request *request)`

Operations on request: `MPI_Wait(MPI_Request request, MPI_Status *status)`; Wait for completion, `MPI_Test(MPI_Request request, int *flag, MPI_Status *status)`; → flag=1 if completed.

MPI I/O: `MPI_File_open(MPI_Comm comm, const char *filename, int amode, MPI_Info info, MPI_File *fh)`; amode → (lecture 10, p. 70), info can be `MPI_INFO_NULL`. `MPI_File_close(MPI_File fh)`; for closing. I/O: Sequential Writing: `MPI_File_write_ordered(MPI_File fh, const void *buf, int count, MPI_Datatype datatype, MPI_Status *status)`; → Writes are ordered by rank. Hybrid MPI Initialization: `MPI_Init_thread(int argc, char **argv, int required, int *provided)` with levels: 1.) `MPI_THREAD_SINGLE` 2.) `MPI_THREAD_FUNNELED` (only "master" makes MPI call(s) 3.) `MPI_THREAD_SERIALIZED` (no overlapping MPI call(s) 4.) `MPI_THREAD_MULTIPLE`. Need to check if provided requested, abort in this case (lecture 10, p. 35)

⑦ **HW Code HW 01 Diffusion Benchmark**: Includes: #include <chrono>, Time measurement: $\text{clock} \text{d}t \text{dt} \{ V[k+d] = ([D-1-k]*D+d]; \}$ reduce Dimensionality: #pragma omp parallel
 for(int i=0; i< Ntimesteps; i++) { const auto t0 = std::chrono::steady_clock::now(); for(int n=0; n< N; n++) { for(int k=0; k< N; k++) { double sum = 0.0; for(int d=0; d< d); uold, unew, N, c); const auto t1 = std::chrono::steady_clock::now(); tsum += t1-t0; std::cout << "swap(" << uold << ", " << unew << ");" } (tsum declared as std::chrono::duration<double> tsum(0)). Performance calculation: const double flop = 5; const double Nflop = flop * N * Ntimesteps; const double fperf = Nflop / tsum.count() * 1.0e-9; **HW 02 Brownian Motion**: Local Generators: std::struct Gen { std::default_random_engine gen; char p[64]; }; std::vector<Gen> gg (omp_get_max_threads()); for(size_t p=0; p< gg.size(); p++) { gg[p].gen.seed(p); } Particle initialization: std::vector<double> xx(N); std::uniform_real_distribution<double> dis(-0.5, 0.5); for(size_t i; i< N; i++) { xx[i] = dis(gg[0].gen); } M random walk steps: #pragma omp parallel { size_t p=omp_get_thread_num(); auto& g = gg[p].gen; std::normal_distribution dis(0, std::sqrt(dt)); #pragma omp for nowait for(size_t i=0; i< Ni; ocnOutputs; o++); { output[0+k*nOutputs] += weights[0*i+nOutputs] * input[i+k*nInputs]; } for(size_t m=0; m< Mi; m++) { xx[i] += dis(g); } } Histogram: std::vector<double> hh(nb, 0); #pragma omp parallel { std::vector<double> (hh(nb, 0)); #pragma omp for nowait int o=0; ocnOutputs; o++); { gradient[0+k*nOutputs] += output[0+k*nOutputs] + (input[i+for(size_t i; i< xx.size(); i++) { int j = (xx[i]-xmin) / (xmax-xmin) * nb; j = std::max(0, std::min(nb-1, j)); hh[j] += 1; } #pragma omp critical for(size_t j=0; j< hh.size(); j=0, sizeof(double)*nOutputs+nInputs); for(int k=0; k< batch_size; k++); { for(int i=0; i< nInputs; i++); hh[j] += (hh[j]); } } **HW 03 Transpose**: for(int d=0; d< D; d++) { for(int n=0; n< N; n++) { uts; i++); for(int o=0; ocnOutputs; o++); { double sum = 0.0; for(int m=0; m< Mi; m++); data[T[d*N+n]] = data[n*D+d]; } } Mean: #pragma omp parallel for for(int d=0; d< D; d++) { m++); sum += output[m+k*nOutputs] + weights[m+i*nOutputs]; } gradient[0*i+nOutputs] = mean[d] = 0; for(int n=0; n< N; n++); { mean[d] = data[T[d*N+n]]; } std::mean[d] = data[T[d*N+n]] / N; } std::+= output[0+k*nOutputs] + (input[i+k*nInputs] - sum); } } compute Eigenvalues: for(int d=0; d< D; d++) { std[d] = 0; for(int n=0; n< N; n++); { double temp = data[T[d*N+n]] for(int k=0; k< batch_size; k++); { for(int o=0; ocnOutputs; o++); { mean[o] += output[0+k*nOutputs] - mean[k*nOutputs]; } } for(int o=0; ocnOutputs; o++); { mean[o] /= batch_size; } for(int o=0; Major: #pragma omp parallel for for(int d=0; d< D; d++); { for(int n=0; n< N; n++); { data[T[0*k*N+i]] = double sum = 0.0; for(int k=0; k< batch_size; k++); { double temp = data[T[d*N+n]] - mean[d]; } std[d] = std::sqrt(std[d] / (N-1)); } standardize Col - mean[k*nOutputs]; } for(int o=0; ocnOutputs; o++); { mean[o] /= batch_size; } for(int o=0; ocnOutputs; o++); { mean[o] = sum / (batch_size-1); } } **HW 04 2D Diffusion**: advance: MPI_Status status[2]; int prev_rank=rank; for(schedule(dynamic, 4) for(int j=0; j< D; j++); { for(int k=0; k< j; k++); { double cov = 0.0; size - 1); } } getEigenvectors: #pragma omp parallel for for(int k=0; k< N; k++); { for(int d=0; next_rank >= procs_); { next_rank = MPI_PROC_NULL; } MPI_Sendrecv(&rho_, for(int i=0; i< N; i++); { cov += data[T[j*N+i]] * data[T[k*N+i]]; } (D+j+k) = cov / (N-1); int next_rank=rank+1; if (prev_rank < 0) { prev_rank = MPI_PROC_NULL; } if (1.0); } } getEigenvalues: #pragma omp parallel for for(int k=0; k< N; k++); { for(int d=0; next_rank >= procs_); { next_rank = MPI_PROC_NULL; } MPI_Sendrecv(&rho_, for(int k=0; k< N; k++); { data_red[n*N+k] = sum; } } reconstructDatasetRowMajor: #pragma omp parallel for for(int n=0; n< N; n++); { for(int d=0; d< d); data_rec[n*D+d] = sum; } } Computing Eigenvalues/Eigenvalues: char job2 = 'V'; char uplo = 'U'; int info, lwork; double * W = new double[D]; double * work = new double[2]; (work[0]=1; work[1]=0; lwork=(int) work[0]; delete[] dsyev_(&job2, &uplo, &D, data_csr, &D, W, work, &lwork, &info); (work[0]=(int) work[0]+1; work[1]=0; lwork=(int) work[0]; delete[] dsyev_(&job2, &uplo, &D, data_csr, &D, W, work, &lwork, &info); (compression Ratio: dim_old = N*d; dim_comp = num_cump + (D+1)*l+2*d); ratio = dim_old / dim_comp; Oja's Rule: forward: for(int k=0; k< batch_size; k++); { for(int i=0; i< nInputs; i++); { for(int o=0; ocnOutputs; o++); { output[0+k*nOutputs] += weights[0*i+nOutputs] * input[i+k*nInputs]; } } ojasRuleGradient: for(int k=0; k< batch_size; k++); { for(int i=0; i< nInputs; i++); { for(int o=0; ocnOutputs; o++); { gradient[0+k*nOutputs] += output[0+k*nOutputs] + (input[i+for(int i=0; i< nInputs; i++); { int j = (xx[i]-xmin) / (xmax-xmin) * nb; j = std::max(0, std::min(nb-1, j)); hh[j] += 1; } #pragma omp critical for(size_t j=0; j< hh.size(); j=0, sizeof(double)*nOutputs+nInputs); for(int k=0; k< batch_size; k++); { for(int i=0; i< nInputs; i++); hh[j] += (hh[j]); } } } sangersRuleGradient: memset(gradient, std::min(int(nb)-1, j)); hh[j] += 1; } #pragma omp critical for(size_t j=0; j< hh.size(); j=0, sizeof(double)*nOutputs+nInputs); for(int k=0; k< batch_size; k++); { for(int i=0; i< nInputs; i++); hh[j] += (hh[j]); } } Covariance: for(int d=0; d< D; d++); { for(int n=0; n< N; n++); { uts; i++); for(int o=0; ocnOutputs; o++); { double sum = 0.0; for(int m=0; m< Mi; m++); data[T[d*N+n]] = data[n*D+d]; } } constructCovariance: #pragma omp parallel output[0+k*nOutputs] - mean[o]; sum += temp * temp; } eigenvalues[o] = sum / (batch_size-1); } } **HW 04 2D Diffusion**: advance: MPI_Status status[2]; int prev_rank=rank; for(schedule(dynamic, 4) for(int j=0; j< D; j++); { for(int k=0; k< j; k++); { double cov = 0.0; size - 1); } } getEigenvectors: #pragma omp parallel for for(int k=0; k< N; k++); { for(int d=0; next_rank >= procs_); { next_rank = MPI_PROC_NULL; } MPI_Sendrecv(&rho_, for(int k=0; k< N; k++); { data_red[n*N+k] = sum; } } reconstructDatasetRowMajor: #pragma omp parallel for for(int n=0; n< N; n++); { for(int d=0; d< d); data_rec[n*D+d] = sum; } } Computing Eigenvalues/Eigenvalues: char job2 = 'V'; char uplo = 'U'; int info, lwork; double * W = new double[D]; double * work = new double[2]; (work[0]=1; work[1]=0; lwork=(int) work[0]; delete[] dsyev_(&job2, &uplo, &D, data_csr, &D, W, work, &lwork, &info); (work[0]=(int) work[0]+1; work[1]=0; lwork=(int) work[0]; delete[] dsyev_(&job2, &uplo, &D, data_csr, &D, W, work, &lwork, &info); (compression Ratio: dim_old = N*d; dim_comp = num_cump + (D+1)*l+2*d); ratio = dim_old / dim_comp;

