

1) Image Processing 1: The Digital Image
 1.1) **Image:** Image is a 2D signal (function depending on some variable with physical meaning) where the input variables are xy-coord. (+ time for videos) and the function value is est. A complete segmentation is a set of regions R_1, \dots, R_N s.t.: $I = \bigcup_{i=1}^N R_i$ and $R_i \cap R_j = \emptyset \forall i \neq j$.
 usually the brightness (or temperature, depth, ...): $f: \Omega^n \rightarrow S$, $n=2$ and $S=\mathbb{R}^+$ for shading: Simple seg. process that produces binary image. $B(x,y)=1$ if $I(x,y) \geq T$.

1.2) Pixel: A pixel is a sample of an original image (not a truelled background ("greenscreen"), then e.g. $I_{\alpha} = I_{\alpha} \cdot g + T$, $g = [0, 255, 0]$). Problems: Variation can be diff (little square)
1.3) Digital Camera: A Charge Coupled Device (CCD) with a sensor array (array of photosites that contain charge proportional to light intensity during exposure), analog to digital converter (measures + digitizes result). happens line by line) **Blooming:** Photosites have finite capacity \rightarrow High intensity leads to "bright" lines" around bright objects:

Bleeding/Smearing: During

$P = TP + FN$, $N = FP + TN$. If V_H, V_P, C_F, C_B are values/costs, choose point on ROC with gradient

2) Image Segmentation
 Image segmentation partitions an image into regions of interest. A complete segmentation is a set of regions R_1, \dots, R_N s.t.: $I = \bigcup_{i=1}^N R_i$ and $R_i \cap R_j = \emptyset \forall i \neq j$.
2.1) Thresholding: Simple seg. process that produces binary image. $B(x,y)=1$ if $I(x,y) \geq T$.
2.2) Chromatotyping: (on greyscale images)

2.3) Performance: True class. needs to be known (by hand). Binary class. has TPs (correctly "in"), TNs (correctly "out") FPs (wrongly "in") and FNs (wrongly "out")
ROC curve: Plots sensitivity ($\frac{TP}{P}$) against 1-specificity ($\frac{FN}{N}$)

3) Pixel Connectivity: 4-neighborhood  8-neighborhood:  Connected Components

4) Labeling: Labels each component of a binary image with a separate number

5) Region Growing: Start from seed/point/region, add neighbour pixels until image edge but not as sharp/lit. **Dark Current:** Small current when no photons are present \rightarrow noise for dark images.

6) CMOS: Each sensor has own amplifier + ADC (+ possibly other components \rightarrow "smart pixels"). Allows random pixel access and cheap, but less sensitive

7) Rolling Shutter: Lines are activated sequentially \rightarrow Can lead to artifacts with moving objects.

8) Aliasing: We store a function's value at many points (sampling) and reconstruct the function. If undersampled \rightarrow indistinguishable from other frequencies. **Nyquist Frequency:** A signal's max. frequency must be smaller than half the sampling frequency of a discrete system to allow exact reconstruction.

9) Quantization: Real valued function will get digital values which is lossy (in contrast to sampling which can be lossless).

10) Image Resolution: No. of pixels (e.g. 1024x1024)

11) Radiometric Resolution: How many bits per pixel

12) Usual Quantization Intervals: Grayscale: 8 bit (256 values), RGB: 8 bit/channel

13) Noise: Common Model: Additive Gaussian Noise. $I(x,y) = f(x,y) + c_n$, with $c_n \sim N(0, \sigma^2)$

14) Signal to Noise Ratio: Index of image quality: $SNR = \frac{F}{\sigma}$, $F = \frac{1}{xy} \sum_{x=1}^X \sum_{y=1}^Y f(x,y)$

15) Color Camera: Prism: Light is separated into RGB using a prism. Filter Mask: Multiple filters are rotated in front of camera.

16) Filter Wheel: Multiple filters are rotated in front of camera.

17) Hit-and-Miss Transform: $H = I \otimes S$. Exact match of S , S can have x ("don't care"). E.g. corner detector:  Thinning: $I \ominus S = I \setminus (I \otimes S)$

18) Skeleton / Medial Axis Transform: Union of centres of maximal discs within a region: 

Calculated: $S = \bigcup_{n=1}^N S_n(x) = (X \otimes_n B) \setminus E(X \otimes_n B) \otimes B$. Skeleton: $S(x) = \bigcup_{n=1}^N S_n(x)$

② 3: Convolution / Filtering

3.1) Filtering: Modifying pixels based on function of local neighborhood in (m, n) / (p, q) space for each edge pixels and increment bins. 3.) Detect peaks in plane

Low Pass: Attenuates ("abschwächen") high frequencies. Specifically \rightarrow smoothing. High Pass: Attenuates low freq. Edge enhancement/detection, suppression of constant gray areas. Band Pass: Attenuates very high approximated by $\Delta^T M \Delta$, $\|M\|=1$, M (structure matrix) = $\begin{bmatrix} \sum_{i,j} w_{i,j} & \sum_{i,j} w_{i,j} \cdot v_{i,j} \\ \sum_{i,j} w_{i,j} \cdot v_{i,j} & \sum_{i,j} w_{i,j} \cdot v_{i,j}^2 \end{bmatrix}$, $w_{i,j} = \frac{1}{\text{neighborhood}}$, $v_{i,j} = \frac{\partial f}{\partial x}$.

Invariant Filter: Linear combination of neighbors, doing same for cell neighbors (shift-invariant). 3.2) Linear Shift-Invariant Filter: Want find $f(x, y)$ where M has large eigenvalues $\rightarrow \lambda_1 \gg \lambda_2$ or $\lambda_2 \gg \lambda_1$ = edge, λ_1, λ_2 and small = "flat region". If $L[[aB(I_1 + B_2)] = aL[I_1] + B[L[I_2]]$, $I'(x, y) = \sum_{(i, j) \in M(x, y)} K(x_i, y_j)$. If K independent $(x, y) \rightarrow$ shift-invariant.

Correlation: $I' = k_0 I$, $I'(x, y) = \sum_{j=i-1}^i \sum_{j'=y-1}^{y+1} K(i, j) I(x+i, y+j)$ [template matching]. 3.3) Correlation: $I'(x, y) = \sum_{i=-3}^3 \sum_{j=-3}^3 K(i, j)$ [point spread function, K is what happens with single pixel]. If $K(i, j) = K(-i, -j) \Rightarrow$ same orientation and scale.

LP Filter: Blur: $\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$ Sharpening: $\frac{1}{8} \begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix} - \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$ Weighted smoothing: $\frac{1}{16} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$ or $\frac{1}{16} \begin{bmatrix} 2 & 2 & 2 \\ 2 & 1 & 2 \\ 2 & 2 & 2 \end{bmatrix}$

3.4) Convolution: $I'(x, y) = \sum_{i=-3}^3 \sum_{j=-3}^3 K(i, j) I(x+i, y+j)$ [template matching]

3.5) Separable Kernels: Can be written as: $K(m, n) = f(m)g(n)$ [Convolution on intermediate result $\rightarrow N^2$ to Fourier transform represents function on new basis with basis elements $e^{-i2\pi(f_m x + g_n y)}$].

2N complexity 3.6) Scale Space: (convolution of Gaussian with itself is Gaussian with sd. $\sqrt{2}$). Repetition convolution with Gaussian \rightarrow Scale space of image.

3.7) Differential Filters: Differential in x-direction $\frac{\partial f}{\partial x}$ approx. by [-1 0 1]. Prewitt: $\begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$. Sobel: $\begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$ (in x-dir., y-dir. by rotating 90°). Used for edge detection.

3.8) Image Sharpening: $I' = I + k(I - I)$, where k is a HP filter.

3.9) Image Features: 4.1)

Template Matching: Locate object, described by template $t(x, y)$ in $s(x, y)$. \rightarrow Minimize MSE: $E(p, q) = \sum_{x=0}^m \sum_{y=0}^n [s(x, y) - t(x-p, y-q)]^2 = \sum_{x=0}^m \sum_{y=0}^n |s(x, y)|^2 + \sum_{x=0}^m \sum_{y=0}^n |t(x-p, y-q)|^2 - 2 \sum_{x=0}^m \sum_{y=0}^n s(x, y) \cdot t(x-p, y-q) \rightarrow$ equiv. to max. "area correlation" = convolution of $s(x, y)$ with $t(x-p, y-q)$.

4.2) Edge Detection: Detect local grad. with finite differences: [-1 0 1], Prewitt, Sobel or Roberts: $\begin{bmatrix} 0 & 1 & 0 \\ -1 & 0 & 1 \end{bmatrix}$. Alternative: Zero-crossings of Laplacian. The FT of the product of two functions is the convolution of their Fourier transforms: $F \cdot G = U(f * g)$.

Detector: 1.) Smoothing with Gaussian 2.) Compute gradient magnitude/angle: $M(x, y) = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$, $\alpha(x, y) = \tan^{-1}\left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x}\right)$ 3.) Nonmaxima suppression to grad. magnitude image (edge normal quantized to 0°, $\pm 45^\circ$, $\pm 90^\circ$; if $M(x, y)$ smaller than any of the two neighbors \rightarrow suppress)

4.3) Hough Transform: 1.) Subdivide (m, n) or (p, q) plane into bins. 2.) Draw line thresholding (strong edge if $M(x, y) \geq \theta_{\text{high}}$, weak if $\theta_{\text{high}} \geq M(x, y) \geq \theta_{\text{low}}$) 5.) Reject weak edges not connected to strong ones.

4.4) Sinusoid Test: $f(x, y) = \sin(\omega_x x + \omega_y y)$. 4.5) Distance from Origin: $d = \sqrt{x^2 + y^2}$. 4.6) Orange: Detect peaks in plane

4.7) Corner Detection: Edges only will be localized in one direction \rightarrow use corners. Displacement sensitivity

4.8) Desired Properties: Accurate localization, invariance against shift/rotation/translation, feature recovery with position, orientation and scale.

4.9) Scale: Strong DoG responses in scale space. Orientation: Gradient directions in patch are calculated, most seen direction is "canonical orientation" of feature. SIFT descriptor: Used for matching, patch is divided into 4x4 regions, for each region canonical orientation is recorded. 4.10) Fourier Transform: Fourier transform represents function on new basis with basis elements $e^{-i2\pi(f_m x + g_n y)}$.

Properties: 1.) Transform is linear 2.) Inverse FT: $f = U^{-1}F$ or $g(x, y) = \iint_{\mathbb{R}^2} F(u, v) e^{i2\pi(f_u x + g_v y)} du dy$ 3.) Scaling: Compress function down \leftrightarrow Stretch function up 4.) FT of Gaussian is a Gaussian

5.1) Examples: Sin wave: $\sum_{x=0}^m \sum_{y=0}^n f(x, y) \leftrightarrow \sum_{x=0}^m \sum_{y=0}^n \frac{1}{\sqrt{2\pi}} e^{-\frac{(x-x_0)^2}{2}} e^{i2\pi f y} \leftrightarrow$ Impulse: $\sum_{x=0}^m \sum_{y=0}^n \delta(x, y) \leftrightarrow \sum_{x=0}^m \sum_{y=0}^n \delta(x-x_0, y-y_0) \leftrightarrow$ Delta: $\sum_{x=0}^m \sum_{y=0}^n \delta(x-x_0, y-y_0) \leftrightarrow$ Box: $\sum_{x=0}^m \sum_{y=0}^n \frac{1}{A} \mathbb{1}_{\{x \in [x_0, x_0+A], y \in [y_0, y_0+A]\}}$

5.2) Examples: Box: $\sum_{x=0}^m \sum_{y=0}^n \frac{1}{A} \mathbb{1}_{\{x \in [x_0, x_0+A], y \in [y_0, y_0+A]\}}$

5.3) Convolution Theorem: The Fourier transform of the convolution of two functions is the product of their Fourier transforms: $F \cdot G = U(f * g)$.

5.4) Sampling: Modeled with δ -function: $\text{Sample}_{2D}(f(x, y)) = f(x, y) \sum_{i=-\infty}^{\infty} \sum_{j=-\infty}^{\infty} \delta(x-i, y-j)$. HF (lead to problems with sampling, should be suppressed before sampling).

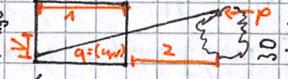
5.5) Computation: 2D FT can be computed as seq. of 1D FTs: $F = \text{Fft}(f)$

5.6) Image Restoration: $s(x) \rightarrow h(x) \rightarrow g(x) \rightarrow f(x) \rightarrow s(x)$, i.e. the "inverse" kernel $\tilde{h}(x)$ should compensate image degradation, $(\tilde{h} \circ h)(x) = \delta(x)$ or $\int \tilde{h}(u, v) f(u, v) du dv = s(x)$

6. Unitary Transforms: Image as a column vector of length $M \times N$, $\vec{f} = A \vec{s}$. Linear image processing is then $\vec{z} = A \vec{f}$. Matrix A (size $M \times N$)

③ is unitary if $A^{-1} = A^H$ (for real $A \Rightarrow A^{-1} = A^T \Rightarrow$ orthogonal). Unitary transforms are energy conserving (conserve vector lengths, are rotation and potentially sign flip of coord. system). G.1) Energy Distribution: $F = [f_1, f_2, \dots, f_n]$ is the image collection, $R_{ff} = E[f_i f_j^H] = \frac{1}{n} F F^H$ (with subtracted means \rightarrow covariance matrix) expected value of the components. Unitary transforms conserve energy, but can distribute it unevenly. $R_{cc} = E[c_i c_j^H] = E[A f_i f_j^H A^H] = A R_{ff} A^H$. Mean squared values ("average energies") of coefficients c_i are of wavelets that are generated by two operations: Scaling (for capturing info at diff. freq.) + translation (diff. location). Le Gall filters: Perfect LP/HP filters would be infinitely large \rightarrow Use 2 analysis/synthesis filters. 7.4) JPEG 2000: 1.) Color components are transformed 2.) Image is tiled (possibly 7 tiles) 3.) Tiles $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_m)$ 3.) R_{ff} is symmetric non-negative definite $\Rightarrow \lambda_i \geq 0 \forall i$ 4.) R_{ff} is a normal matrix: are wavelet transformed 4.) Quantization 5.) Coding. Steps 7, 3, 4 can be lossy/lossless 8.) Optical Flow: $R_{ff}^H R_{ff} = R_{ff}^H R_{ff}$, unitary eigenmatrix therefore exists. The KL-transform is a unitary transform with Try to track where pixels move to. Used for tracking, structure from motion, motion segmentation, stabilization, etc. 8.1) Brightness Constancy: OF is the apparent motion of brightness patterns. Let $A = \phi^H$, i.e. $\tilde{c} = \phi^H f$, where the values of ϕ are ordered according to decreasing eigenvalues. Transform, normalization, compression, ... conservation: No other unitary transform puts as much energy into the first j coeff., the MSE by choosing only first j coeff. is minimized. Proof: Consider $\tilde{\delta}^j = \sum_i \tilde{c}_i \delta_i^j$. $E = \text{Tr}(P_{bb}) = \text{Tr}(\sum_i \tilde{c}_i \delta_i^H A R_{ff} A^H) = \sum_i \tilde{c}_i^2 \delta_i^H R_{ff} \delta_i$. Introducing Lagrange multipliers $\sum_k \lambda_k (1 - \tilde{c}_k^T \delta_k)$ and differentiating w.r.t. \tilde{c}_j gives: $\text{Residual} = \delta_j$ as Eigenimages: for $I_x = \frac{\partial I}{\partial x}$, $I_y = \frac{\partial I}{\partial y}$, $I_z = \frac{\partial I}{\partial z}$ any unitary transform, the inv. transform $\tilde{f} = A^H \tilde{c}$ can be interpreted as "superposition" of "basis images" to $I(x, y, z) = I(x, y, z) + \frac{\partial I}{\partial x} \delta_1 + \frac{\partial I}{\partial y} \delta_2 + \frac{\partial I}{\partial z} \delta_3$. Leads (with the assumption) to the constraint equation: $\frac{\partial I}{\partial x} \delta_1 + \frac{\partial I}{\partial y} \delta_2 + \frac{\partial I}{\partial z} \delta_3 = 0$. 8.2) Aperture Problem: Let $u = \frac{dx}{dt}$, $v = \frac{dy}{dt}$, $I_x = \frac{\partial I}{\partial x}$, $I_y = \frac{\partial I}{\partial y}$, $I_z = \frac{\partial I}{\partial z}$. Min. $e_s + \alpha e_c$ [iter. methods, finite differences]. 8.4) Lukas-Kanade Algorithm: Single velocity for all pixels of patch. $E(u, v) = \sum_{(x,y) \in \text{patch}}$ Subtract mean face 3.) Calculate $p = \phi_K^H x$, compare with DB entries. 6.3) LDA: LDA maximizes between-class variance, while minimizing within-class variance. Leads to "fishertaces" \rightarrow better with varying lightning. 7.) Pyramids/Wavelets 7.7) JPEG: Humans don't resolve high freq. too well \rightarrow used by JPEG, applying discrete cosine transform (DCT) on 8x8 blocks. DCT coeffs. are quantized (potentially lossy) and ordered according to freq. ("zig-zag" pattern ). First coeff. = DC, only diff. to prev. DC encoded. Sequential \rightarrow Use different levels of image pyramid to estimate at different scales. 8.6) Parametric Motion Models: Use not only translation, but also affine/quadratic models. E.g. affine: $E(A, b) = \sum [(I(Ax) - I_b)^2]$ Mann encoding (least bits for most common words) used. 7.2) Image Pyramid: Image blurred and subsampled Offer more constrained solutions than smoothness/integration over larger areas than translation-only. at every level. Used to search for correspondence, edge tracking, computational cost. Gaussian pyramid \rightarrow Redundant because LP filter. Laplacian Pyramid: Diff. between upsampled Gaussian and Gaussian Laplacian \rightarrow Brightness constancy with noise $I(x, y, t) = I(x + u \delta t, y + v \delta t, t + \delta t) + \eta$, η ind. Gaussians. 8.7) Bayesian Optic Flow: Can explain some low-level human motion illusions by adding uncertainty to

④ Arrive at $v^* = -\begin{bmatrix} \sum I_x^2 + \frac{C}{\sigma^2} \\ \sum I_x I_y \\ \sum I_y^2 + \frac{C}{\sigma^2} \end{bmatrix}^{-1} \begin{pmatrix} \sum I_x I_t \\ \sum I_y I_t \end{pmatrix}$ → One param: Ratio of observation/prior spread. transforms of $u(v, r)$. $R(p, \theta)$ is called sinogram because RT of off-center point is sinusoid

8.8) SSD: Can use template matching for large displacements. 9) Video Compression: Visual Per 7.1) Backprojection: Given $R(p, \theta)$, can we find $u(r, v)$? → ill-posed problem, backprojection projection (limited to <24Hz, but still need to avoid aliasing). Interlaced video: Two half images (eventually) project measured sinograms back into image space to construct $u(r, v)$. 7.2) Central Slice Theorem: rows) displayed consecutively, increasing frequency. Full image repr. = Progressive. Lossy video compres. The FTs of $R(p, \theta)[\tilde{f}(v, \theta)]$ and $u(r, v)[\tilde{u}(r, v, \theta)]$ are related. For a given θ , $\tilde{f}(v, \theta) =$ sion takes advantage of redundancy: Spatial correlation between neighboring pixels, temporal corr. $\tilde{u}(r, v, \theta, r, v, \theta)$. Therefore, for all θ : 1.) Measure attenuation $R(p, \theta)$ 2.) 1D FT: $\tilde{f}(v, \theta)$ 3.) 2D IFT ation between frames. 9.1) Temporal Redundancy: Current frame is predicted based on prev. coded fr. and sum with prev. values to reconstruct $u(r, v)$. After step 2, HP filter for noise reduction comes. 3 types of frames: 1.) I-frame: Intra-coded frame coded independently 2.) P-frame: Predictively-coded. motion can be applied ("filtered backprojection"). Computer Graphics Use of computers to synth- ed frame based on prev. frame 3.) B-frame: Bidirectionally predicted frame, based on prev. and future frame. For P/B-frames, only residual is encoded. Better performance with motion compen- sation (MC): Motion described on block-basis with motion vector. Vector generated by exhaustive search We have $\frac{v}{z} = \frac{y}{z}$, therefore $v = \frac{y}{z}$ and likewise $u = \frac{x}{z}$ 12.1) Perspective Projection: Where does a point $p=(x, y, z)$ end up in the image? Assuming pinhole model, unit size camera: 

can be based on prev./future, both (aug.) or none. GOP: MPEG has group of pictures; say, that st- re "virtual diamond" is intersected:  12.2) Rasterization: Conversion of a continuous or "multistep" approach: Coarse grid search+refinement (up to subpixels). For B-frames, MC-prediction object to a discrete representation on a pixel grid. Line Rasterization: Only light pixels up who- avgs with I-frames and ends with frame before next I-frame, and has defined nr. of P/B- els, for a line with points $(u_1, v_1), (u_2, v_2)$, slope is $s = \frac{v_2 - v_1}{u_2 - u_1}$. Travel to next pixel according frames → Allows decoding without going to start of video. 9.2) Spatial Redundancy: Frames/residuals are encoded similarly to images (DCT) 9.3) PSNR: Error for 1 pixel: Diff. between original/decoded $e_{mse} = \sqrt{\frac{1}{N \cdot M} \sum_{i=1}^N \sum_{j=1}^{M-1} (x_{ij} - \hat{x}_{ij})^2}$. PSNR: $\frac{\text{max. value of } X}{e_{mse}}$ 10) Sparsity and Texture: 10.1) Image Denoising: Try to re- pixel according to fraction of pixel area covered by triangle, which can be estimated by samp- present image as sparse vector of basis images (designed, e.g. wavelets, or learned) to reduce noise. 10.2) Lighting (multiple measurements per pixel). Rule needed when edge is on sampling point (OpenGL): Representing Texture: Want to find the stylized subelements of texture and represent their stats. To find them, filters are applied and response magnitude is recorded. 10.3) Texture Synthesis: Image can be used as source of prob. model. Better results by also modeling co-occurrence. Alternative: Image-based approaches ("copy-pasting"), e.g. chesses mosaic that tiles image, picks random blocks and smooths edges. Or choose blocks s.t. overlap is maximized.

11) Radon Transform In computer tomography, we can measure reduction in intensity of an X-ray, given by $I = I_0 e^{-\mu \int L u(s) ds}$, where $u(s)$ is the optical density of the material (what we want to estimate).

Line L can be reparametrized: $L = (p \cos(\theta) - s \sin(\theta), p \sin(\theta) + s \cos(\theta))$, which gives $R(p, \theta) = \int_{-\infty}^{\infty} u(s) \delta(p - s \cos(\theta) - s \sin(\theta)) ds = \int_{-\infty}^{\infty} u(r, v) \delta(p - r \cos(\theta) - v \sin(\theta)) dr dv$, which is the radon

transforms of $u(r, v)$. $R(p, \theta)$ is called sinogram because RT of off-center point is sinusoid

12.2) Perspective Projection: Given $R(p, \theta)$, can we find $u(r, v)$? → ill-posed problem, backprojection pro- 3.) Composition of linear transforms is linear → uniform representation. Key idea is that Her take lines to lines while keeping the origin fixed. 14) Trans- 13.1) Triangle Coverage: (Can color 13.2) Triangle Coverage: Point-in-triangle test is used to test if sample is inside. Given a triangle $P_i = (X_i, Y_i)$ ($i \in \{0, 1, 2\}$), $dX_i = X_{i+1} - X_i$, $dY_i = Y_{i+1} - Y_i$, we define $E_i(r, v) = (x - X_i)dY_i - (y - Y_i)dX_i$. Point is inside it $E_0(r, v) < 0, E_1(r, v) < 0, E_2(r, v) < 0$. Instead of testing all points, various traversal orders possible → GPUs often do tiled triangle traversal with parallel testing in blocks. 14) Trans- 14.1) Linear Maps: Reasons: 1.) Computationally easy to solve eq. 2.) Over short distance, all rays approx. by linear one 3.) Composition of linear transforms is linear → uniform rep-

⑤ Map f is linear if it maps vectors to vectors and $f(\mathbf{a}+\mathbf{b})=f(\mathbf{a})+f(\mathbf{b})$, $f(c\mathbf{a})=c f(\mathbf{a})$. Linear maps between \mathbb{R}^m and \mathbb{R}^n can be expressed as: $f(\mathbf{u}) = \sum_{i=1}^n u_i \mathbf{a}_i$ (for fixed vectors \mathbf{a}_i). For \mathbb{R}^2 , $f(\mathbf{u}) = u_1 \mathbf{a}_1 + u_2 \mathbf{a}_2$ with $u_1 = f(e_1), u_2 = f(e_2)$, i.e. we only need to know $f(e_1), f(e_2)$ to map the entire space.

14.2) Scale: Uniform Scale: $S(x) = ax$, Non-uniform: $S(x) = x_1 a_1 + x_2 b_2$, $A = \begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix}$

14.3) Rotation: $R_\theta =$ rotate counter-clock by θ . Rigid/isometric \rightarrow shape (distance between stored, surface is where interpolated values = 0)

Reflection: $R_y(x) =$ reflection about y -axis. $R_y(x) = -x_1 e_1 + x_2 e_2$, $A = \begin{bmatrix} -1 & 0 \\ 0 & 1 \end{bmatrix}$

14.4) Shear: $H(x) = x_1 \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} + x_2 \begin{bmatrix} 0 & 1 \\ 0 & 0 \end{bmatrix}$ at all scales. L-system: Set of grammar rules, start symbol, semantics. E.g. grammar $\{A \rightarrow BBA, B \rightarrow abA\}$

14.5) Shear: $H(x) = x_1 \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} + x_2 \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$ with start symbol AAA. Semantics could be a = forward 10, b = left hand(10,50), ... (can generate e.g. fractals)

14.6) Translation: Not linear, but affine $[f(x) = ax+b]$ is affine, not linear

14.7) Matrix: Can write linear transformations as matrix-vector product: $\begin{bmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = x_1 \begin{bmatrix} a_{11} \\ a_{21} \end{bmatrix} + x_2 \begin{bmatrix} a_{12} \\ a_{22} \end{bmatrix}$

14.8) 2D Homogeneous Coordinates: 2D-point (x_1, x_2) is represented as 3D vector $\begin{bmatrix} x_1 \\ x_2 \\ 1 \end{bmatrix}$. E.g. rotation becomes with very simple rules.

14.9) Composition: By matrix mult. Order right-to-left, $f(x) = R \circ S \circ x$ is scaling, then rotation. Rotate about point x : 1.) Translate by $-x$ 2.) Rotate 3.) Translate by x

14.10) 3D-H: E.g. shear (in x , based on y, z): $\begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} = \begin{bmatrix} x_1 \\ x_2 \\ x_3 + b x_1 \end{bmatrix}$

14.11) 3D rotations: About x -axis: $\begin{bmatrix} 1 & 0 & 0 \\ 0 & \cos\theta & -\sin\theta \\ 0 & \sin\theta & \cos\theta \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}$, y -axis: $\begin{bmatrix} \cos\theta & 0 & \sin\theta \\ 0 & 1 & 0 \\ \sin\theta & 0 & \cos\theta \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}$, z -axis: $\begin{bmatrix} \cos\theta & -\sin\theta & 0 \\ \sin\theta & \cos\theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}$

14.12) Camera Transform: Want to compute transforms from world to camera coord. system (camera at origin, looking down $-z$). 1.) Translation by $-x$ (x =camera pos.) 2.) Form orthonormal basis V around camera direction, (u, v, w) , triangles as triplets of indices (i, j, k)

14.13) Perspective Projection: Given $\frac{x}{w}, \frac{y}{w}, \frac{z}{w}$ let Θ be the field of view in y -direction, $f = \cot(\frac{\Theta}{2})$ and ϕ angle between light direction/surface normal. Given flux Φ and $\text{dot}(N, L)$

14.14) View Frustum: Region in space that will appear on screen: $P(0,0) = \frac{f}{r}, P(1,1) = f, P(2,2) = \frac{f^2}{r}, P(2,3) = \frac{f^2 + 2f\text{near} + 2\text{far}}{2\text{near} + 2\text{far}}$

14.15) Screen Transform: After projection, all points from $(-1, -1)$ to $(1, 1)$ to $\max(0, \text{dot}(N, L))$. common abstraction: infinitely bright light source „at infinity“ \rightarrow same L .

14.16) Color: $(1, 1)$ are on screen, \rightarrow need to transform to $(0, 0)$ to (w, h) for output: 1.) Reflect about x -axis 2.) fr -When assigning colors to triangle vertices \rightarrow barycentric coord. for in-between colors. Given triangle (a, b, c) , translate by $(1, 1)$ 3.) Scale by $(w/2, h/2)$

15: Geometry

15.1) Implicit Representations: Points aren't known directly, but satisfy some relationship. Pros: 1.) Can be compact 2.) Easy inside/outside check 3.) Other queries (e.g. distance) may also be easy 4.) Exact description 5.) Easy topology changes. Cons: 1.)

15.2) Explicit Representations: Points are given directly \rightarrow easy sampling, hard in-side/outside tests. Point Cloud: List of points (x, y, z) , sometimes augmented with normals. Drawing in undersampled regions can be difficult. Polygonal Mesh: Definition polygon: Vertices V_0, \dots, V_{n-1} , edges E : set of connected polygons $M = \langle V, E, F \rangle$ (planar + non self-interacting). Polygonal mesh: Set of connected polygons

15.3) Geometry Sources: 3D scanning (point cloud common), digital 3D modeling (\rightarrow subdivision surfaces).

16: Texture

16.1) Lambert's Law: Irradiance at surface prop. to $\cos\theta$ $\text{dot}(N, L)$. Dot product between surface normal (N) and unit direction to light (L), i.e. $\text{dot}(N, L) = \frac{\text{dot}(N, L)}{\|L\|}$

16.2) Color: $X_{color} = aX_{color} + bX_{color} + cX_{color}$

16.3) Barycentric Interpolation: Can also be computed as area ratios: $a = \frac{A_a}{A}, b = \frac{A_b}{A}, c = \frac{A_c}{A}$. Barycentric interpolation in screen space: 1.) Compute depth z at each vertex

6) vertex. 2.) Eval. $Z = \frac{1}{2}, P = \frac{f}{2}$ 3.) Interpolate Z/P with barycentric coord. 4.) Divide interpolated P by interpolated Z to get f. **16.3) Texture Mapping:** Many uses: Surface material prop., normal mapping, texture value used to perturb surface normal(), precomputed lightning/shadows. Texture coord. define mapping from surface coord. (points on triangle) to points in texture domain. They are linearly interpolated over the triangle → only defined for vertices. Texture mapping (for each covered screen sample)

- 1.) $(u,v) = \text{eval. texcoord. at } (r,r)$
- 2.) Sample texture at (u,v)
- 3.) Set samples color to tex. color.

Texture sample positions may not be uniformly distributed in texture space. To avoid aliasing → pre-filter textures. Mipmaps = Precomputed texture at different resolutions/filter levels. (choose level according to texture coord. distance of neighboring pixels (far objects → strongly filtered level). **17.1) Occlusion:** Depth buffer (2-buffer) can be used to decide which object is visible at each pixel. Stores the depth of closest triangle that has been processed so far. Processing order doesn't matter, works with interpenetrating triangles/supersampling.

Rasterization Pipeline

- 1.) Occlusion: Depth buffer (2-buffer) can be used to decide which object is visible at each pixel. Stores the depth of closest triangle that has been processed so far. Processing order doesn't matter, works with interpenetrating triangles/supersampling.
- 2.) Compositing: Alpha = additional image channel for opacity ($\alpha=1 \rightarrow$ fully opaque, $\alpha=0 \rightarrow$ fully transparent). "over" operator composites image A over B. $A \text{ over } B \neq B \text{ over } A$. If we composite B over A, composited color is: $C = \alpha_B A + (1-\alpha_B) \alpha_A A$. (can also store $A' = [\alpha_A A_r \alpha_A A_g \alpha_A A_b \alpha_A]$, then $C = B' + (1-\alpha_B) A'$). Rendering: Check if depth test passed, update color but not depth buffer (no direction).
- 3.) Render opaque surfaces using normal 2-buffer
- 4.) Disable 2-buffer updates, render semi-transparent objects

17.2) Compositing: "over" operator composites image A over B. $A \text{ over } B \neq B \text{ over } A$. If we composite B over A, composited color is: $C = \alpha_B A + (1-\alpha_B) \alpha_A A$. (can also store $A' = [\alpha_A A_r \alpha_A A_g \alpha_A A_b \alpha_A]$, then $C = B' + (1-\alpha_B) A'$). Rendering: Check if depth test passed, update color but not depth buffer (no direction).

17.3) Real-Time Gouraud Shading Pipeline: Input: Vertices in 3D space. 1.) Vertex operations (transformation matrices) 2.) OpenGL's (law). Fragment operations: Fragment generation (rasterization) which produces fragment streams, fragment processing (with textures) which takes streams and produces shaded fragment stream 4.) Operations on screen samples (depth/color).

Programs: Define behavior of vertex/fragment processing stages. Describe operation on single vertex/fragment.

18.1) Light: Light is electromagnetic radiation, color determined by the frequency. "White" light = mixture of all visible frequencies. Emission spectrum of a light source (additive model): How much light at different wavelengths is produced → useful for comparing light sources. Absorption spectrum (subtractive model): How much light is absorbed → we

source (additive model): How much light at different wavelengths is produced → useful for comparison/characterizing color of point, ink, ... Brain: We have S, M, L cones (short, medium, long wavelengths) in eyes, brain interprets (S, M, L) response as color. **18.2) Color Spaces:** Color from some space (like an artist's palette) is specified using a color model. Additive models e.g. used for combining colored lights) like RGB, subtractive models (e.g. used for combining print colors) like CMYK. There's also HSV(hue, saturation, value) and YCbCr(intensity, blue-yellow deviation from gray, red-cyan deviation). **18.3) Radiometry:** Assume photons travel in straight lines. Each photon carries energy, want to record amount of energy (brightness) **Radiant Energy:** Total number of photons hitting scene. **Radiant Flux:** Hits per second **Irradiance:** Hits per second per unit area. **Radiance:** Energy along ray, defined by origin point p and direction w . Complete description of light in an environment → rendering is about computing radiance. Need to distinguish between incident ($L_i(p,w)$) and exitant ($L_o(p,w)$) radiance at point on surface. In general, $L_i(p,w) \neq L_o(p,w)$.

18.4) Isotropic Point Source: Slightly more realistic lighting model, where light spreads in every direction. Irradiance from radiance: $E = \int_{H^2} L_i(p,w) \cos \theta dw$, where θ is angle between w and normal.

18.5) Rendering Equation: $L_o(p,w_0) = L_e(p,w_0) + \int_{H^2} f_r(p,w_i \rightarrow w_0) L_i(p,w_i) \cos \theta_i dw_i$, where: L_o = outgoing radiance, L_e = emitted radiance (from e.g. light source), f_r = reflection/scattering function and describes how reflection affects outgoing radiance. Choice determines surface appearance: 1.) Ideal specular (perfect mirror) 2.) Ideal diffuse (uniform reflection in all directions) 3.) Glossy specular (majority reflected in reflection direction) 4.) Retro-reflective (back towards source)

18.6) BRDF: Bi-directional reflectance distribution function, type of reflection function with properties: 1.) $f_r(w_i \rightarrow w_0) \geq 0$ 2.) $\int_{H^2} f_r(w_i \rightarrow w_0) \cos \theta_i dw_i \leq 1$ 3.) Helmholtz-reciprocity: $f_r(w_i \rightarrow w_0) = f_r(w_0 \rightarrow w_i)$ [violated by subsurface scattering]

19. Raytracing: We "shoot" a ray for

⑦ every screen sample, color sample according to hit object. Proceed in screen sample order (vs. raster event size → useful for non-uniform primitive distribution). Unlike BVH, can terminate search after first hit.

ization, which proceeds in triangle order). No depth buffer needed and natural order for prims. Quad-trees (2D) / octrees (3D): Uniform grid, but cells can be further divided (by 4/8) if they have too many transparent surfaces (under they appear along ray), but entire scene must be stored. Can view many primitives. **18. Animation** (creating illusion of motion, e.g. by rapidly displaying minimally changing images).

rasterization as efficient algorithm for visibility queries, given rays with same origin and uniform. Some principles: 1.) Squash and stretch (shape distortion) 2.) Anticipation (action preparation) 3.) Follow-through distribution over projection plane. **19.1) Shadows:** For rasterization → shadow mapping: Scene is rendered (termination of action) 4.) Slow in, slow out (timing). Hand-drawn motion: Keyframes by senior artist, in-between (depth buffer only) from light location, this buffer is then used when rendering from camera POV even by constant. Techniques in computer animation: 1.) Artist-directed (e.g. keyframing) 2.) Data-driven (e.g. for shadow queries). For raytracing → "shadow rays" are shot towards light source from place where motion capture) 3.) Procedural (e.g. simulation) **20.1) Keyframing:** Specify important events (position, color, lights, camera rays intersect scene. **19.2) Reflections:** For rasterization → Render scene 6 times (all directions) from camera configuration... only, interpolate next. **Splines:** Any piecewise polynomial: 1D: (t_i, f_i) , $f(t_i) = f_i$. Has to reflect object's POV, apply renderings as textures to object. For raytracing → Shoot another ray when hitting reflective object (can model surfaces by reflection function) → recursive raytracing.

19.3) Ray Intersection: Ray given by $r(t) = o + td$, o = origin, d = dir. vector. Intersection with implicit surface: Plug in r for x , solve for t . E.g. plane, given by $N^T x = c \rightarrow N^T(o + td) = c \rightarrow t = \frac{c - N^T o}{N^T d}$, i.e. $o + \frac{c - N^T o}{N^T d}$ as intersection point. Ray-triangle intersection: Intersect with plane, compute $\frac{\text{spurture inside}}{\text{spurture outside}}$ to check for intersection point. **19.4) Ray-Scene Intersection:** Given scene with M primitives, find first hit object. To do better than $O(N)$ → Put axis-aligned boxes around objects, only consider primitives in hit boxes. **Bounding Volume Hierarchy:** Divides space into set of bounding boxes, organized tree-like. Interior nodes represent subset of primitives, leaf nodes contain list of primitives. Primitives are partitioned into disjoint sets, but boxes can overlap in space. **Front-to-back** traversal works like this: 1.) Check if box not hit or closest point on box farther than curr. min. distance → return 2.) If leaf node → Traverse all primitives 3.) Otherwise, traverse first child where closest point is nearer. To implement partitioning → Can use axis-aligned approach. Split plane/line along axis is chosen, primitives are partitioned according to the side of their centroid. → Can evaluate small number of split planes with a heuristic to get plane. **Space-Partitioning Acceleration Structures:** In contrast to BVH, can also partition space instead primitives into disjoint sets → Primitive can be contained in multiple regions. E.g. uniform grid where each cell contains primitives that overlap it → Allows efficient walking through similar to line rasterization). K-D trees partition space via axis-aligned planes (partitions can have overlaps).

⑧ Animation (creating illusion of motion, e.g. by rapidly displaying minimally changing images.)

camera configuration... only, interpolate next. **Splines:** Any piecewise polynomial: 1D: (t_i, f_i) , $f(t_i) = f_i$. Has to be polynomial between knots: $t_i \leq t \leq t_{i+1}$, $f(t) = \sum_{j=0}^3 c_i t^{i+j} =: p_i(t)$. Most common $cl-3$ (min. number of 2nd derivatives for small displacements). Runge's phenomenon: Higher degree can lead to worse approx. (oscillation). **Natural Splines:** p_i : cubic. 1.) Interpolation: $p_i(t_i) = f_i$, $p_i(t_{i+1}) = f_{i+1}$, $i = 0, \dots, n-1$ 2.) C^1 continuity: $p_i'(t_{i+1}) = p_{i+1}'(t_{i+1})$ 3.) C^2 continuity: $p_i''(t_{i+1}) = p_{i+1}''(t_{i+1})$. 4.) DCFs, only $n=2$ eq. $\rightarrow p_0''(t_0) = p_n''(t_n) = 0$ **Spline Properties:** 1.) Interpolation 2.) C^2 continuity 3.) Locality (moving one control point doesn't affect whole curve). Not all possible at the same time (natural: no locality) **Hermite/Bézier Splines:** Tangents are also specified (won't have sent is calculated as diff. of neighbors: $U_i := \frac{f_{i+2} - f_{i-1}}{t_{i+2} - t_{i-1}}$) **B-Spline:** No. (exact) interpolation. $B_{i,k}(t) = 1$ for $t_i \leq t \leq t_{i+k}$, 0 otherwise. $B_{i,k}(t) := \frac{t - t_i}{t_{i+k} - t_i} B_{i,k-1}(t) + \frac{t_{i+k} - t}{t_{i+k} - t_{i-1}} B_{i-1,k-1}(t)$ $f(t) = \sum_{i=0}^n a_i B_{i,k}(t)$ **Rigging:** Character animation pipeline = Modeling, rigging, animation. Rig = User-defined mapping between params and deformations of high-res mesh. Rigging = Creating a rig. Animation then defined through set of (spline) curves, determining rig params over time. **21.1) Blend Shape Rigs:** Input: Meshes M_i with vertices x_i^j , blending weights $a_i^j = (a_{i,1}, \dots, a_{i,n})$. Output: Blended mesh M through linear combination: $M = \sum_i a_i M_i$, $x^j = \sum_i a_i^j x_i^j$. Keyframes: Blending weights $a_i(t_i)$. **21.2) Cage-based Deformers:** High-res model embedded in a coarse mesh (cage), cage is deformed and deformations are applied to high-res model. For a point x_i in the high-res model: $x_i = \sum_j w_i^j c_j$; (c_j = points of cage, w_i^j = weights). **21.3) Skeletal Animation:** Animate only skeleton, simulate muscles, tissue, skin, environment interactions... (computationally expensive). **21.4) Linear Blend Skinning:** Deformations of surface mesh are coupled to skeletal pose. **Forward Kinematics:** Joints = Local coord. frames,

(8) Bones: vectors between consecutive joints. Each non-root bone is defined in frame of unique parent, whose transformation to parent frame affect all descendant bones. $p(j)$: Parent of joint j ; frame of joint j expressed w.r.t. its parent: $p(j)R_j = \begin{pmatrix} R_{p(j)} & t(j) \\ 0 & 1 \end{pmatrix}$, i.e. rotation + translation. $t(j)$ typically constant during animation, $R_{p(j)}$ varies. Transformation from frame j to world: $wR_j = wR_0 \dots p(j)R_{p(j)}p(j)R_j = t(0)R_{p(0)} \dots t(p(j))R_{p(j)}t(j)R_{p(j)}$

Skinning: Given joint angles (i.e. rotation matrices), skeleton conf. can be computed using FK. How to adjust surface mesh? Rigid skinning: Assign each vertex to closest bone, compute world coord. according to bone's transformation: $V = wR_j w\bar{R}_j^{-1} V'$, V' = vertex coord. in rest pose (first transformed in coord. frame of bone j , so wR_j can be applied). Works fine if vertices close to bone/small deformations. Linear blend skinning: Vertex assigned to multiple bones with weights (manually painted or computed) or Neumanns (derivative is specified).

inverse Kinematics: Given goal position for "end effector" (e.g. hand), compute joint angles to achieve position. \rightarrow Analytic formulation, Energy-based methods, ... **27.5 Motion Capture**: We get joints target by MC. Becomes optimization-based III: $\min_{\theta} f_0(\theta) = \frac{1}{2} (\rho(\theta) - \rho)^T (\rho(\theta) - \rho)$, ρ = target, $\rho(\theta)$ = MC-transformed point. Side constraints (e.g. max. joint angle) possible.

22. Physics-Based Animation: Newton's 2nd law: $F = ma$, F = forces, m = mass, a = acceleration. More general: System has config. $q(t)$, velocity $\dot{q} = \frac{d}{dt} q$, mass M , forces F are acting on it and potentially constraints $g(q, \dot{q}, t) = 0$. Then, $\ddot{q} = \frac{F}{m}$, want to solve for $q(t)$.

(Complex phenomena modeled as collection of particles with attached forces, acting on them 22.7 Numerical Integration: Given $q(0), \dot{q}(0)$, find $q(t)$. Function $q(t)$ replaced by discrete samples q_i : $q_{i+1} = q_i + \Delta q_i$, $\Delta q_i = \frac{q(t+h) - q(t)}{h}$ approx. $\frac{1}{2}(f(t) + f(t+h))$.

1.) Rectangle rule: $\Delta q_i = q(t) \cdot h$ 2.) Midpoint rule: $\Delta q_i = q(t + \frac{h}{2}) \cdot h$ 3.) Trapezoid rule: $\Delta q_i = \frac{q(t) + q(t+h)}{2} \cdot h$ **Forward Euler**: $q_{i+1} = q_i + h \dot{q}_i \rightarrow$ Only stable for small timesteps **Backward Euler**: $q_{i+1} = q_i + T f(q_{i+1}) \rightarrow$ Unconditionally stable for linear ODEs. **Symplectic Euler**: Only for 2nd order, no numerical damping/energy is preserved. Velocity updated using curr. config., config. using new velocity. **23. PDEs**: $u_t = \frac{\partial^2 u}{\partial x^2}$

$u_{xx} = \frac{\partial^2 u}{\partial x^2}, \nabla = \left(\frac{\partial}{\partial x_1}, \dots, \frac{\partial}{\partial x_n} \right)^T, \Delta = \nabla^T \cdot \nabla = \frac{\partial^2}{\partial x_1^2} + \dots + \frac{\partial^2}{\partial x_n^2}$

23.7 (Conservation): Order: order of highest partial derivative, linear function u and its derivatives only occur linearly (e.g. $u_t + c \cdot u_x = 0$ vs. non-linear $u_t + c \cdot u_x = 0$) \rightarrow coeffs can be non-linear functions. 2nd order linear PDEs have form: $A u_{xx} + 2B u_{xy} + C u_{yy}$

(wave eq.) $u_{yy} = F(x, y, u, u_x, u_y)$. 1.) **Hyperbolic**: $B^2 - A > 0$ (wave front is hard to solve/no steady state) 2.) **Parabolic**: $B^2 - A \leq 0 \rightarrow$ Heat eq. (spread of heat over time) solution same as Laplace after long time) 3.) **Elliptic**: $B^2 - A < 0 \rightarrow$ Laplace eq. (smoothest function interpolating boundary data) easy to solve robust to errors)

23.2 Numerical Solution: Pick spatial + time discretization, run time-stepping. Spatial = Lagrangian/Particles (conceptually easy, no resolution/domain limitation, good distribution can be hard, finding neighbors can be expensive) or Eulerian/Grid (fast, regular comp., easy representation, "trapped" in grid) **Laplace Eq.**: $\frac{u_{i+1,j} - 4u_{i,j} + u_{i-1,j} - u_{i,j+1} + u_{i,j-1}}{h^2} = 0$. Solve $au = 0$ by weighted avg. Need to approx. Laplacian, e.g. $\frac{1}{4}u_{i+1,j} - \frac{4}{4}u_{i,j} + \frac{1}{4}u_{i-1,j} - \frac{1}{4}u_{i,j+1} + \frac{1}{4}u_{i,j-1}$.

Heat Eq.: $u_t = \Delta u$. E.g. forward Euler, $u^{k+1} = u^k + \Delta u^k$, where Laplacian is approx. as above. **Wave Eq.**: $u_t = \Delta u \rightarrow u = v, v = \Delta u$ **Exercises**: BG subtraction: Exemplar set model: Use set of BG pixels, calc. mean μ / cov. covariance Σ . BG if $\text{sort}((\mu-\mu)^T \Sigma (\mu-\mu)) \gg 1$. Per-pixel model: For each pixel, mean/cov. matrix calculated over BG frames. Then, class classification rule used. **Filtering**: Creation of Gaussian HP: Extract HP content by subtracting LP filtered img. from original \rightarrow HP filter constructed by subtracting LP from unit response. **PCA Face Detection**: Lowest compression error \rightarrow most likely face location. **Transformations**: Rigid-body transformation: "motor-translation" \rightarrow rotation + translation. Rotations commute only in 2D. **Coord. system change**: Given local coord. system with axes x, y, z and origin t , $F = [x \ y \ z]^T$ $P_{world} = F_p$, if p is expressed in coord. system. **PVM**: $M = [x \ y \ z \ t]$ is model matrix that expresses model coord. system. C is the camera matrix, typically $C = TR$ (translation/rotation) $P_{cam} = M_p C_p \rightarrow p_{cam} = C^{-1} M_p p$. V is the view matrix, $V = C^{-1} \cdot R^{-1} \cdot T^{-1} = \tilde{R}^T$. **Screen projection** from camera point with projection matrix P , therefore $P_{cam} = P V M_p p$. **Phong Shading**: Phong reflection model: $I = I_{amb} + I_{diff}(N \cdot L) + I_{spec}(R \cdot V)^n$, I : ambient, I_{diff} : diffuse, I_{spec} : specular, K : material constant, K_d/K_s : reflection factor, n : exponent for surface character (smooth \Rightarrow 32, perfect mirror \Rightarrow 1). L : light dir., N : surface normal, V : viewing dir., R : reflection dir. with $R = 2(L \cdot N)N - L$. **Linear Interpolation**: $y_0 + (t-t_0) \# (y_1-y_0)/(t_1-t_0)$

Appendix: Assumptions Lucas Kanade: 1.) Brightness constancy 2.) Small motion 3.) Spatial coherence: All points in neighborhood have same motion. **Radians**: $SI = 720^\circ$ **Trigonometry**: $\sin(0^\circ) = 0, \cos(0^\circ) = 1, \sin(90^\circ) = 1, \sin(-90^\circ) = -1, \cos(-90^\circ) = 0$ **Table of Contents**: 1: The Digital Image: p. 1, 2: Image Segmentation: p. 1, 3: Convolution / Filtering: p. 2, 4: Image Features: p. 2, 5: Fourier Transform: p. 2, 6: Unitary Transforms: p. 2/3, 7: Pyramids / Wavelets: p. 3, 8: Optical Flow: p. 3, 9: Video Compression: p. 4, 10: Scarcity and Textures: p. 4, 11: Radon Transform: p. 4, 12: Drawing Triangles: p. 4, 13: Transform: p. 4, 14: Geometry: p. 5, 16: Texture: p. 5, 17: Rasterization Pipeline: p. 6, 18: Rendering Equation: p. 6, 19: Raytracing: p. 6/7, 20: Animation: p. 7, 21: Rigging: p. 7, 22: Physics-Based Animation: p. 8 **Appendix II: Fourier Transform Sampling Visualization**

10 Fourier Transform:

No Aliasing: signal \xrightarrow{FT} magnitude $\xrightarrow{\text{convolve}}$ sample $\xrightarrow{\text{mult.}}$ aliasing! **Aliasing**: signal \xrightarrow{FT} magnitude $\xrightarrow{\text{copy/shift}}$ sample $\xrightarrow{\text{mult. with box filter}}$ aliasing! **Matrix Multiplication**: $[a \ b \ c] \cdot [g \ j] = [ag+bi+cj] \dots$ $[d \ e \ f] \cdot [h \ k] = [dh+ei+fj] \dots$ $[2 \times 3] \quad [3 \times 2] \quad [2 \times 2]$